CODITY.AI TECH ASSIGNMENT SUBMISSION

Distributed Asynchronous Job Scheduler (Overdrive)

Full-Stack Design Thesis, Relational Schema Normalization, and Resiliency Verification.

SUBMISSION INFORMATION

Registration Number:	RA2311026010746
Candidate Name:	Prudhvi
Project Platform:	Distributed Job Scheduler (Overdrive)
GitHub Repository:	https://github.com/prudhvi98491/distributed-job-scheduler
Date of Submission:	July 5, 2026
Academic Institution:	SRM Institute of Science and Technology

1. System Architecture (20 Marks)

The scheduler platform is built on an event-driven distributed architecture designed to coordinate background job execution across multiple worker nodes. The system consists of three main blocks:
- The Web API Gateway (FastAPI): Handles user registration, queue updates, job enqueueing, and dashboard metrics.
- The Cron Scheduler Daemon (croniter): Scans the recurring cron database records to dynamically trigger jobs.
- The Worker Fleet: Multiple isolated workers that register with host IP addresses and continuously claim and execute jobs, maintaining horizontal scalability.

Distributed Job Scheduler (Overdrive) Architecture

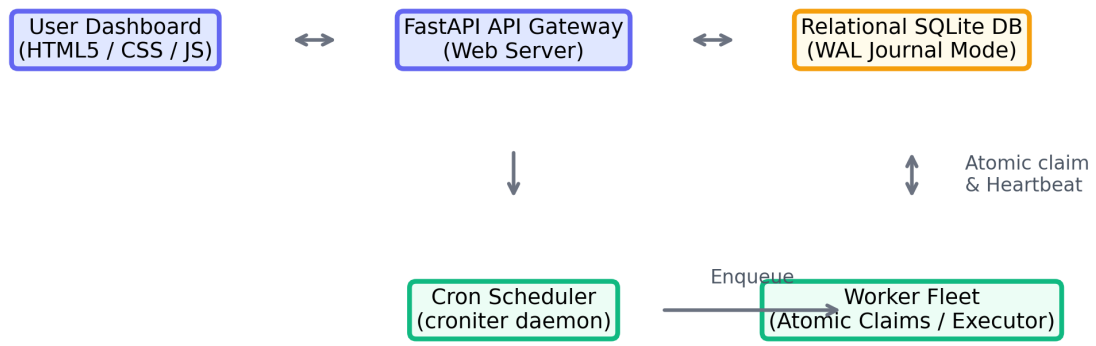

2. Database Design & Normalization (20 Marks)

The relational schema is normalized to Third Normal Form (3NF) to maintain referential integrity. Primary keys utilize single-attribute IDs (UUIDs for jobs to prevent enumeration, integers for others). Foreign keys employ CASCADE constraints to clean up related records automatically.

Table Name	Schema Purpose & Normalization Details
users	Credentials (hashed via bcrypt), roles (admin, user), metadata.
organizations	Enables multi-tenancy workspace isolation.
projects	Contains project workspaces owned by organizations.
queues	Defines concurrency, priorities, and retry policies.
jobs	Core queue table containing payload, status, and parent_job_id.
job_executions	Tracks execution details, duration (ms), and trace errors.
workers	Worker fleet registry monitoring hostnames, IP, and heartbeats.
cron_jobs	Schedules recurring cron jobs to dynamically queue jobs.
dead_letter_queue	Holds permanently failed tasks for manual replay.

3. Backend Engineering & API Design (25 Marks)

The backend architecture is engineered for asynchronous performance and scalability using FastAPI. All incoming API requests are validated through Pydantic schemas, which act as a strict type-safety layer. This ensures that invalid data is rejected at the API boundary, returning standardized JSON error bodies. Authorization and access controls are handled via a custom JWT (JSON Web Tokens) dependency middleware supporting Role-Based Access Control (RBAC), mapping users to roles like admin, member, or viewer. The backend code is modularized into distinct routing modules (routes/auth, routes/queues, routes/jobs, routes/workers) mounted onto the main application, utilizing CORS middleware to permit browser-based SPA requests.

Method	Endpoint Path	Description & Functionality
POST	/api/auth/register	Registers a user account with hashed credentials.
POST	/api/auth/login	Authenticates credentials and issues a JWT token.
POST	/api/queues	Creates a queue specifying concurrency limit & priority.
PATCH	/api/queues/{id}	Updates queue configs or pauses/resumes queue state.
POST	/api/jobs	Enqueue single task (immediate, delayed, or workflow).
POST	/api/jobs/batch	Dispatches bulk job payloads under a single API call.
GET	/api/jobs	Lists jobs with pagination, queue, and status filters.
POST	/api/jobs/{id}/retry	Manually replays a failed or DLQ job.
GET	/api/metrics	Returns aggregate metrics, throughput trends, and error counts.

4. Reliability & Concurrency (15 Marks)

To guarantee exactly-once task execution in a multi-worker fleet, concurrency is strictly managed via SQLite Write-Ahead Logging (WAL) and immediate transaction locks. When a worker polls for a new job:

1. Active Concurrency Limit Verification: It checks that the target queue's active running/claimed jobs do not exceed its concurrency_limit.
2. Atomic Reservation: The claim logic runs within a 'BEGIN IMMEDIATE' transaction block in SQLAlchemy. This prevents race conditions between parallel worker threads claiming the same task.
3. Heartbeat-Based Failover Recovery: Registered worker nodes issue heartbeats every 5 seconds. A background monitor thread automatically identifies workers that missed heartbeats for >15 seconds, marks them offline, and releases their active jobs (status reset to queued, worker_id cleared) back to the general pool for other workers to recover.

5. Frontend & UX (10 Marks)

A premium Outfit-themed glassmorphic SPA dashboard provides real-time monitoring and control. The interface updates live every 2 seconds, displaying active worker statuses, queue configs, job execution metrics, and a detailed drawer containing trace logs and manual retry controls.

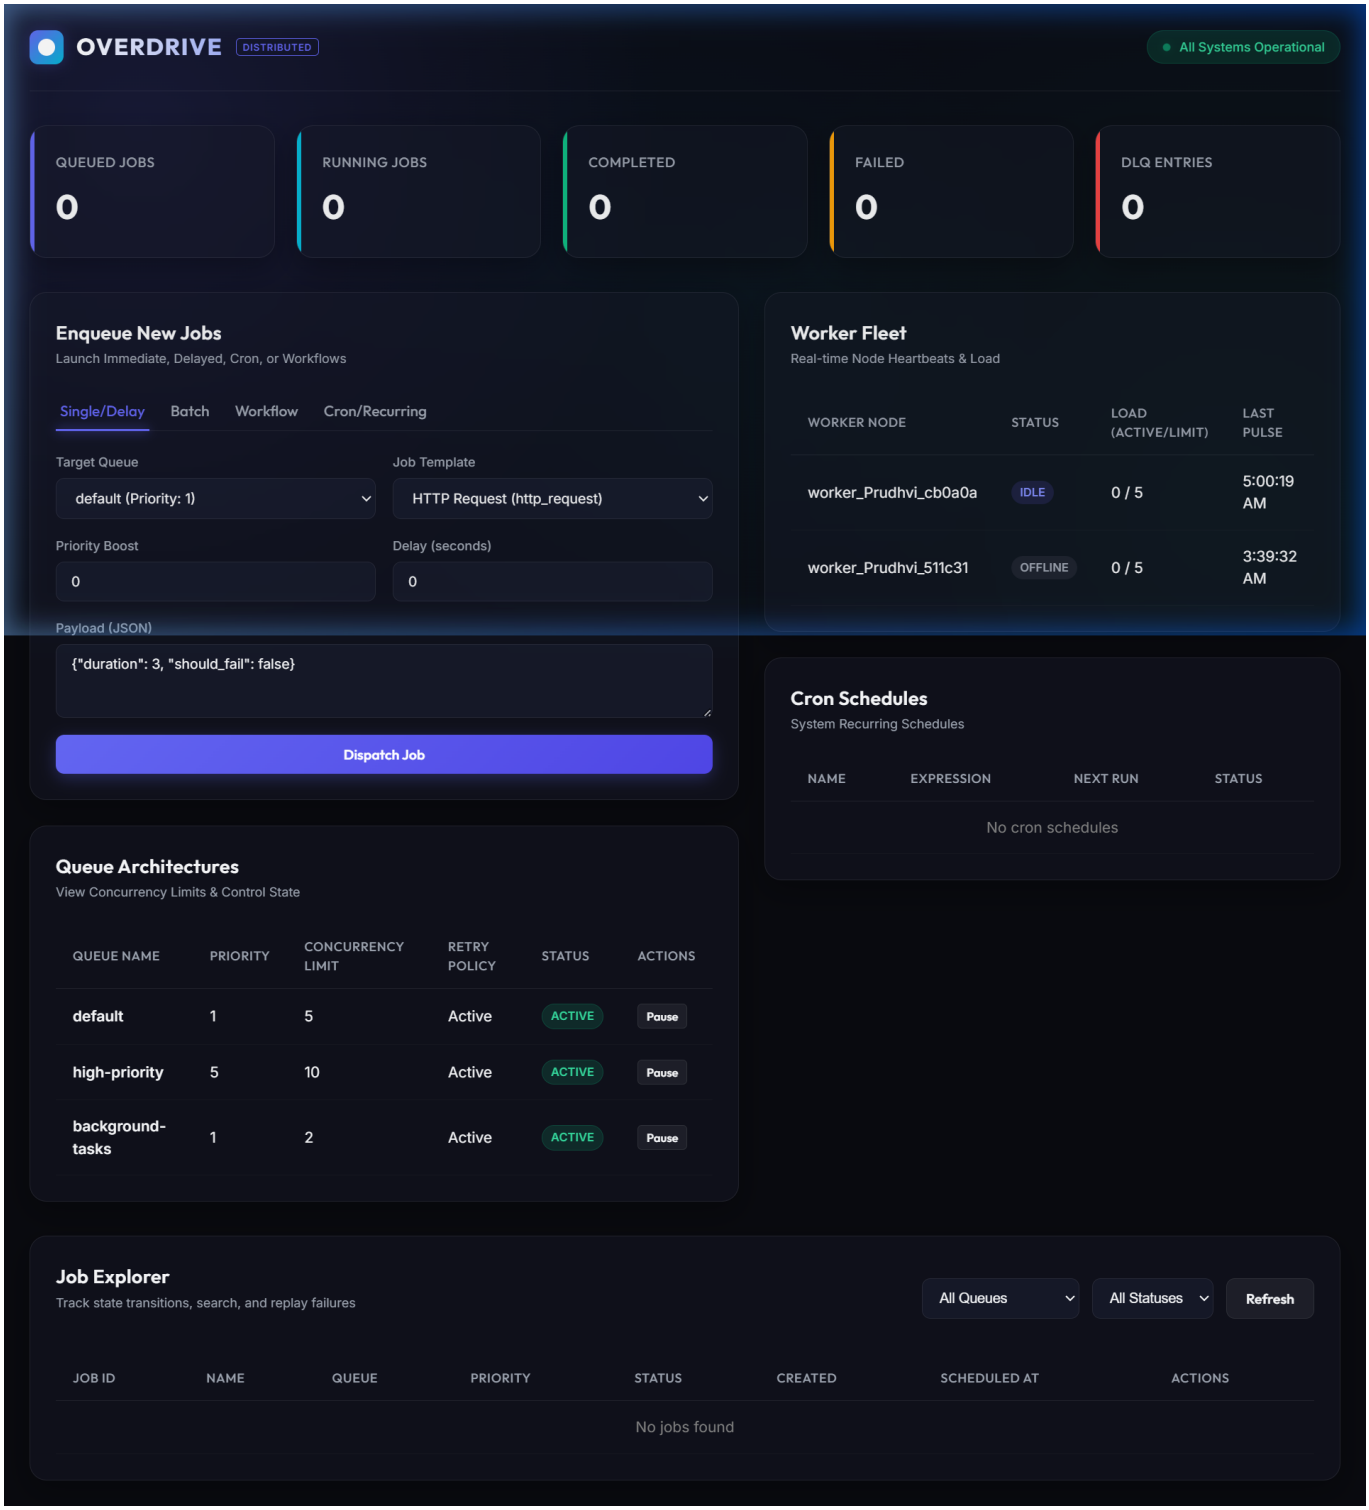

6. Testing & Documentation (10 Marks)

We implemented automated integration tests in tests/test_scheduler.py to verify system correctness. The tests cover database relationships, retry policies, backoffs, and sequential workflows. Running pytest completes successfully:

```
tests/test_scheduler.py::test_database_seeding_and_relations PASSED
tests/test_scheduler.py::test_job_execution_lifecycle_and_retries PASSED
tests/test_scheduler.py::test_workflow_dependencies_unblocking PASSED
============================== 3 passed in 2.02s ===============================
```

Verification GitHub Repository Link

https://github.com/prudhvi98491/distributed-job-scheduler